

Protecting Liquidity Providers Across Time-Regime Discontinuities

A Comprehensive Guide for Tokenized Securities AMMs

Mathematical Frameworks, Risk Models, and Implementation Strategies

StockShield Protocol Research Document

Ayush Popat

January 29, 2026

Contents

1	Introduction: The LP Protection Problem	3
1.1	Why Tokenized Securities Are Different	3
1.2	The Three Risk Surfaces	3
1.3	What This Document Covers	4
2	Part 1: Foundational Academic Models	4
2.1	The Avellaneda-Stoikov Model (2008)	4
2.1.1	The Problem Setup	4
2.1.2	Key Assumptions	4
2.1.3	The Reservation Price	5
2.1.4	The Optimal Spread	5
2.2	Kyle's Model (1985): Price Impact and Informed Trading	6
2.2.1	The Setup	6
2.2.2	Kyle's Lambda	6
2.3	Glosten-Milgrom (1985): The Bid-Ask Spread as Adverse Selection Cost	7
2.3.1	The Logic	7
2.3.2	The Formula	7
2.3.3	Spread Component Decomposition	8
3	Part 2: Understanding LP Losses	8
3.1	Loss-Versus-Rebalancing (LVR)	8
3.1.1	The Intuition	8
3.1.2	The LVR Formula	8
3.1.3	LVR vs. Impermanent Loss	9
3.2	Impermanent Loss: The Detailed Formula	9
4	Real-World Edge Cases: Stories That Illustrate The Risks	10
4.1	Edge Case 1: The Overnight Earnings Surprise	10
4.2	Edge Case 2: The Monday Morning Gap	11
4.3	Edge Case 3: The Halted Stock	11
4.4	Edge Case 4: The Low-Liquidity Hour Attack	12
4.5	Edge Case 5: The Flash Crash Recovery	13
4.6	Edge Case 6: The Informed Trader	13
4.7	Edge Case 7: The Holiday Weekend Disaster	14
4.8	Edge Case 8: The Stale Oracle Attack	15
4.9	Summary: The Common Threads	16
5	Part 3: Risk Detection Mechanisms	16
5.1	Time-Based Regime Detection	16
5.1.1	Deterministic Market Hours Calendar	16
5.1.2	Pseudocode for Regime Detection	16
5.2	VPIN: Volume-Synchronized Probability of Informed Trading	17
5.2.1	The Intuition	17
5.2.2	Calculation Steps	17
5.2.3	Interpreting VPIN	18
5.3	Price Deviation from Oracle	18
5.4	Oracle Staleness Check	18

6 Part 4: Protection Mechanisms	19
6.1 Dynamic Fee Adjustment	19
6.2 Gap Capture Mechanism	19
6.2.1 The Problem	20
6.2.2 The Solution: Minimum Bid with Decay	20
6.2.3 How It Works	20
6.3 Inventory-Aware Quote Skewing	20
6.4 Circuit Breakers	21
6.4.1 Trigger Conditions	21
6.4.2 Graduated Response	21
7 Part 5: Implementation Architecture	21
7.1 System Components	22
7.2 Key Data Structures	22
7.3 Hook Implementation Outline	22
7.4 Volatility Tracking (On-Chain)	23
8 Part 6: Parameter Calibration	24
8.1 Fee Parameters	24
8.2 Circuit Breaker Thresholds	25
8.3 Gap Auction Parameters	25
9 Part 7: Hackathon Prize Track Integration	25
9.1 Technology-to-Purpose Mapping	25
9.2 Uniswap v4: Privacy-Enhancing LP Protection	25
9.2.1 How StockShield Addresses Privacy DeFi Requirements	26
9.2.2 Uniswap v4 Hook Implementation Points	26
9.3 Yellow Network: Off-Chain Scaling via State Channels	26
9.3.1 Why State Channels for LP Protection?	27
9.3.2 Integration Architecture	27
9.3.3 ERC-7824 State Channel Flow	27
9.4 ENS: Decentralized Identity Layer	28
9.4.1 ENS Use Cases in StockShield	28
9.4.2 ENS Implementation	28
9.4.3 Reputation-Based Fee Tiers	29
9.5 Integrated System: Full Stack Architecture	30
9.5.1 Complete Transaction Flow	30
9.6 Prize Track Qualification Matrix	31
10 Conclusion	31

1 Introduction: The LP Protection Problem

Liquidity providers (LPs) in automated market makers (AMMs) face a fundamental challenge: they must continuously offer to buy and sell assets at quoted prices, even when market conditions change unfavorably. In traditional DeFi, LPs accept this risk in exchange for trading fees. However, when AMMs are used for **tokenized securities**—assets that represent ownership in real-world stocks, bonds, or other regulated instruments—the problem becomes significantly more complex.

1.1 Why Tokenized Securities Are Different

Unlike cryptocurrencies that trade 24/7, tokenized securities are linked to underlying assets that trade on traditional exchanges with specific hours:

- **NYSE Regular Hours:** 9:30 AM – 4:00 PM Eastern Time
- **Pre-Market:** 4:00 AM – 9:30 AM Eastern Time (limited liquidity)
- **After-Hours:** 4:00 PM – 8:00 PM Eastern Time (limited liquidity)
- **Overnight:** 8:00 PM – 4:00 AM Eastern Time (no trading)
- **Weekends:** No trading from Friday 4:00 PM to Monday 4:00 AM

This creates a fundamental mismatch: a blockchain-based AMM can technically accept trades 24/7, but the LP has no way to hedge their risk when the underlying market is closed. If news breaks overnight that causes Apple stock to gap down 10% at the next market open, an LP who sold Apple tokens at the old price overnight has suffered a significant loss with no ability to offset it.

1.2 The Three Risk Surfaces

This document addresses three interconnected risk surfaces that all stem from one core issue: **LPs are forced to quote across time regimes where hedging ability, information quality, and market participation change discontinuously.**

1. **Traditional NYSE Open/Close Edges:** Predictable gaps occur at market open and close. News accumulates overnight, and the opening price often differs significantly from the previous close. LPs who quoted during closed hours are exposed to “stale price” arbitrage.
2. **Extended-Hours Transitions:** With proposals for 22-24 hour trading (NYSE Arca, 24Exchange), there are multiple session handoffs throughout the day. Each transition creates mini-gaps and liquidity discontinuities.
3. **Low-Liquidity Hours:** The most dangerous risk surface. During overnight or pre-market hours, one-sided order flow accumulates, there’s no venue to hedge, and informed traders can exploit thin liquidity with minimal cost.

Key Insight

These are not three separate problems but three manifestations of one core issue: **discontinuous hedging ability**. The solution must be a unified framework that detects regimes, adjusts pricing dynamically, and manages inventory proactively.

1.3 What This Document Covers

This guide provides a comprehensive treatment of LP protection, including:

- **Part 1: Foundational Academic Models** – The mathematical frameworks that underpin market making
- **Part 2: Understanding LP Losses** – LVR, impermanent loss, and why LPs lose money
- **Part 3: Risk Detection Mechanisms** – How to detect dangerous market conditions
- **Part 4: Protection Mechanisms** – Dynamic fees, circuit breakers, and auction design
- **Part 5: Implementation Architecture** – Smart contract design for Uniswap v4 hooks
- **Part 6: Parameter Calibration** – How to set the numbers that make it work

2 Part 1: Foundational Academic Models

Before designing protection mechanisms, we must understand *why* LPs lose money and *how* market makers have traditionally managed these risks. This section covers the academic models that form the theoretical foundation.

2.1 The Avellaneda-Stoikov Model (2008)

The Avellaneda-Stoikov model is the foundational framework for optimal market making under inventory risk. It answers the question: “*Given that I’m a market maker with some current inventory, what prices should I quote?*”

2.1.1 The Problem Setup

A market maker faces two opposing forces:

1. **Earning the Spread:** By quoting a bid price below and an ask price above the “true” price, the market maker profits from the difference when trades occur on both sides.
2. **Inventory Risk:** If more buyers than sellers arrive, the market maker accumulates a long position. If the price then falls, they lose money on that inventory. The reverse happens with a short position.

The market maker wants to maximize expected utility of their final wealth, accounting for the risk of adverse inventory positions.

2.1.2 Key Assumptions

- The “true” mid-price of the asset follows a random walk (Brownian motion): $dS_t = \sigma dW_t$
- Buy and sell orders arrive randomly according to Poisson processes
- The arrival rate of orders depends on how far the quote is from the mid-price
- The market maker has a risk aversion parameter γ that penalizes holding inventory

2.1.3 The Reservation Price

The key insight of the model is the concept of a **reservation price**—the price at which the market maker is indifferent to trading. This is **not** the market mid-price, but a personal valuation that accounts for inventory.

Reservation Price Formula:

$$r(s, q, t) = s - q \cdot \gamma \cdot \sigma^2 \cdot (T - t)$$

Where:

- s = current market mid-price
- q = current inventory position (positive = long, negative = short)
- γ = risk aversion parameter (higher = more risk-averse)
- σ = volatility of the asset
- $T - t$ = time remaining until the end of the trading session

How to interpret this formula:

- If the market maker is **long** ($q > 0$), the reservation price is **below** the mid-price. This encourages selling (to reduce inventory) by making the ask more attractive.
- If the market maker is **short** ($q < 0$), the reservation price is **above** the mid-price. This encourages buying (to cover the short) by making the bid more attractive.
- The adjustment is **larger** when:
 - Inventory $|q|$ is larger (more risk exposure)
 - Risk aversion γ is higher (more concerned about risk)
 - Volatility σ is higher (more potential for adverse moves)
 - More time remains ($T - t$) (more time for bad things to happen)

Example

Suppose the mid-price of ETH is \$2,000, the market maker is long 5 ETH, risk aversion $\gamma = 0.01$, volatility $\sigma = 0.5$ (50% annualized), and there are 4 hours left until market close ($(T - t) = 4/8760$ years ≈ 0.000457).

$$r = 2000 - 5 \times 0.01 \times 0.5^2 \times 0.000457 \times 8760 = 2000 - 5 \times 0.01 \times 0.25 \times 4 = 2000 - 0.05 = 1999.95$$

The market maker's personal valuation is \$0.05 below the market mid-price, reflecting the desire to reduce their long position.

2.1.4 The Optimal Spread

Once the reservation price is established, the market maker sets bid and ask quotes at optimal distances from it. The Guéant-Lehalle-Fernandez-Tapia extension (2013) provides the formula:

Optimal Half-Spread:

$$\delta^* = \gamma\sigma^2(T - t) + \frac{1}{\gamma} \ln \left(1 + \frac{\gamma}{\kappa} \right)$$

Where κ is the “order arrival intensity parameter” measuring how sensitive order flow is to price.

The total spread (bid to ask) is $2\delta^*$, so:

$$\text{Bid} = r - \delta^* \quad \text{and} \quad \text{Ask} = r + \delta^*$$

Key insight for AMMs: The spread should **widen** when:

- Volatility increases (more risk per trade)
- More time remains until close (more exposure time)
- Risk aversion increases (preference for safety)

2.2 Kyle's Model (1985): Price Impact and Informed Trading

While Avellaneda-Stoikov focuses on inventory risk, Kyle's model addresses **adverse selection**—the risk of trading against someone who knows more than you do.

2.2.1 The Setup

In Kyle's model, there are three types of market participants:

1. **Informed Trader:** Knows the true value V of the asset and trades strategically to profit from this knowledge.
2. **Noise Traders:** Trade randomly for reasons unrelated to the asset's value (liquidity needs, rebalancing, etc.). Their aggregate order flow is $u \sim N(0, \sigma_u^2)$.
3. **Market Makers:** Observe total order flow $y = x + u$ (informed + noise) but cannot distinguish between them. They set prices to break even on average.

2.2.2 Kyle's Lambda

The key result is that market makers set price as a linear function of order flow:

Price Impact:

$$p = \mu + \lambda \cdot y$$

Where:

- μ = unconditional expected value of the asset
- λ = “Kyle's lambda” = price impact per unit of net order flow
- y = total signed order flow (positive = net buying, negative = net selling)

In equilibrium, Kyle's lambda equals:

$$\lambda = \frac{\sigma_V}{2\sigma_u}$$

Where σ_V is the standard deviation of the informed trader's private information and σ_u is the standard deviation of noise trading.

Interpretation:

- Higher λ means each unit of order flow moves the price more (less liquid market)
- λ increases when informed traders have better information (σ_V high)
- λ decreases when there's more noise trading (σ_u high) to "hide" among

Key Insight

For LP protection, λ is a measure of **toxicity**. When λ is high, order flow is more likely to be informed, and LPs should widen spreads or reduce depth to protect themselves.

2.3 Glosten-Milgrom (1985): The Bid-Ask Spread as Adverse Selection Cost

The Glosten-Milgrom model explains why bid-ask spreads exist even when market makers are risk-neutral and have no operating costs.

2.3.1 The Logic

1. Some traders are **informed**—they know whether the asset is overvalued or undervalued. They only buy when it's undervalued and only sell when it's overvalued.
2. Some traders are **uninformed**—they trade for liquidity reasons regardless of value.
3. The market maker cannot tell who is informed. But they know that:
 - Buy orders are more likely to come from informed traders when the asset is undervalued
 - Sell orders are more likely to come from informed traders when the asset is overvalued
4. Therefore, the market maker adjusts prices:
 - **Ask price** = Expected value *given that someone wants to buy*
 - **Bid price** = Expected value *given that someone wants to sell*

2.3.2 The Formula

If a fraction μ of traders are informed, and the asset can be worth either V_H (high) or V_L (low):

Spread from Adverse Selection:

$$\text{Spread} = \text{Ask} - \text{Bid} = \mu \cdot (V_H - V_L)$$

The spread is proportional to:

- The **fraction of informed traders** (μ)—more informed traders = wider spread
- The **information advantage** ($V_H - V_L$)—bigger potential mispricing = wider spread

Warning

This model has a critical implication: the spread exists **purely to compensate for losing to informed traders**. Even a risk-neutral market maker with zero operating costs must charge a spread, or they will systematically lose money to informed traders.

2.3.3 Spread Component Decomposition

Research by Stoll (1989) decomposed actual bid-ask spreads into their components:

Component	% of Spread	Description
Adverse Selection	43%	Cost of trading with informed traders
Order Processing	47%	Operational costs of market making
Inventory Holding	10%	Capital cost and price risk of holding inventory

For AMM-based LPs, the “order processing” component is essentially zero (smart contracts don’t need salaries), but adverse selection and inventory costs remain.

3 Part 2: Understanding LP Losses

Now that we understand the theoretical frameworks, let’s examine exactly how LPs lose money in practice.

3.1 Loss-Versus-Rebalancing (LVR)

LVR, introduced by Milionis et al. (2023), is the definitive measure of LP costs in AMMs. It captures the **adverse selection cost** that arbitrageurs extract from LPs.

3.1.1 The Intuition

Imagine you’re an LP in a ETH/USDC pool. External markets (like centralized exchanges) update prices continuously based on new information. But your AMM’s prices only update when trades happen. This creates a window where your AMM offers a “stale” price.

Example:

1. At 10:00 AM, ETH trades at \$2,000 on both Coinbase and your AMM.
2. At 10:01 AM, news breaks and ETH jumps to \$2,050 on Coinbase.
3. At 10:01 AM, your AMM still offers to sell ETH at around \$2,000.
4. An arbitrageur buys ETH from your AMM at \$2,000 and sells it on Coinbase at \$2,050.
5. The arbitrageur profits \$50 per ETH. You, the LP, lose \$50 per ETH.

This happens continuously. Every time external prices move, arbitrageurs extract value from the AMM.

3.1.2 The LVR Formula

For a constant-product AMM, the expected LVR rate is elegantly simple:

Instantaneous LVR Rate:

$$\text{LVR Rate} = \frac{\sigma^2}{8}$$

Where σ is the annualized volatility of the asset price.

Example calculations:

Volatility (σ)	Daily LVR	Annual LVR
25%	0.021%	0.78%
50%	0.086%	3.13%
75%	0.193%	7.03%
100%	0.342%	12.5%

Key Insight

LVR is the cost of being a **passive** price setter. The $\sigma^2/8$ formula connects AMM LP returns to options pricing—LPs are effectively short gamma, meaning they lose money when prices move in either direction.

3.1.3 LVR vs. Impermanent Loss

People often confuse LVR with impermanent loss (IL). They are related but distinct:

Property	LVR	Impermanent Loss
Benchmark	Rebalancing portfolio	Holding portfolio
Path-dependent?	Yes	No (only start/end prices)
Reversible?	No (permanently lost)	Yes (if prices return)
Cause	Arbitrage extraction	Price divergence
Formula	$\sigma^2/8$ (continuous)	$2\sqrt{r}/(1+r) - 1$

The key difference: IL measures what you lost compared to just holding the initial assets. LVR measures what you lost compared to a portfolio that continuously rebalanced to match the AMM's composition at fair prices (with no arbitrageur taking the other side).

Warning

Impermanent loss can reverse if prices return to their starting point. LVR is **permanent**—once an arbitrageur extracts value, it's gone forever. LVR is the true cost of being an LP.

3.2 Impermanent Loss: The Detailed Formula

For completeness, here's the impermanent loss formula for a constant-product AMM:

Impermanent Loss:

Let $r = P_1/P_0$ be the ratio of final price to initial price. Then:

$$\text{IL}(r) = \frac{2\sqrt{r}}{1+r} - 1$$

Impermanent Loss Table:

Price Change	r	IL
-50%	0.5	-5.72%
-25%	0.75	-1.31%
0%	1.0	0%
+25%	1.25	-0.62%
+50%	1.5	-2.02%
+100% (2x)	2.0	-5.72%
+200% (3x)	3.0	-13.4%
+400% (5x)	5.0	-25.5%

Notice that IL is always negative or zero—LPs never *gain* from price movements. They only avoid losses when prices don’t change.

4 Real-World Edge Cases: Stories That Illustrate The Risks

The formulas above can feel abstract. This section presents concrete scenarios—stories with real numbers—that show exactly how LPs can get hurt. Each scenario is something that has happened or could easily happen in tokenized securities markets.

4.1 Edge Case 1: The Overnight Earnings Surprise

The Setup: Sarah is an LP providing \$100,000 of liquidity in an Apple/USDC pool on a tokenized securities AMM. It’s Thursday evening, and Apple is trading at \$150 per share. Her pool holds the equivalent of 333 tokenized Apple shares and 50,000 USDC.

What Happens:

1. **Thursday 4:00 PM ET:** NYSE closes. Apple is at \$150. Sarah’s pool is balanced.
2. **Thursday 5:30 PM ET:** Apple releases earnings. They beat expectations massively. In after-hours trading (which Sarah’s AMM can’t hedge against), Apple jumps to \$165.
3. **Thursday 6:00 PM ET:** An arbitrageur sees that Sarah’s pool still offers Apple at \$150. They buy 100 tokens from the pool at an average price of \$152 (there’s some slippage).
4. **The arbitrageur’s profit:** They paid \$15,200 for tokens worth \$16,500, netting \$1,300.
5. **Sarah’s loss:** Her pool gave away \$1,300 in value. She now has fewer Apple tokens (they went to the arbitrageur) and the USDC she received is less than what those tokens are actually worth.

Why This Is Dangerous:

- Sarah had **no way to hedge**. The NYSE is closed. There’s no liquid market where she could buy Apple to offset her pool’s sale.
- The AMM’s oracle might not have updated yet (oracles often lag during after-hours).
- This will repeat whenever news moves prices while traditional markets are closed.

Key Insight

The Protection: During after-hours and overnight periods, the AMM should either pause trading entirely, or charge fees high enough (e.g., 2-5%) to compensate LPs for the inability to hedge. A “gap auction” mechanism can also capture some of this value back for LPs.

4.2 Edge Case 2: The Monday Morning Gap

The Setup: Bob is an LP in a Tesla/USDC pool with \$50,000 of liquidity. On Friday at market close, Tesla is trading at \$200.

What Happens:

1. **Friday 4:00 PM:** NYSE closes. Tesla at \$200. Bob's pool is balanced.
2. **Saturday:** Elon Musk tweets something controversial. Crypto markets (open 24/7) start pricing in the impact. However, Bob's tokenized Tesla cannot trade on NYSE.
3. **Sunday:** More news develops. Market consensus forms that Tesla will open down 15% on Monday.
4. **Monday 9:30 AM:** NYSE opens. Tesla opens at \$170. But for the first few seconds, Bob's AMM still has the pool priced around \$200 until the oracle updates.
5. **The Attack:** High-frequency arbitrageurs immediately sell Tesla tokens into Bob's pool at \$198, then buy them back on NYSE at \$170. They extract \$28 per share.
6. **Bob's loss:** On 50 tokens sold into his pool, that's \$1,400 lost in seconds.

The Numbers:

Item	Value
Gap size	-15% (\$200 → \$170)
Exploitable value per share	\$28 (allowing for slippage)
Shares sold into pool	50
Total extracted by arbitrageur	\$1,400
Bob's remaining pool value	\$48,600

Warning

The Worst Part: This happens in milliseconds. By the time Bob even knows the market has opened, the damage is done. Human reaction time cannot protect against this.

The Protection: At market open, implement a “soft open” period where:

- Trading is paused for 30-60 seconds while oracles update
- Or, trades require a “gap auction bid” that goes directly to LPs
- Or, spreads are set extremely wide (10-20%) until price stabilizes

4.3 Edge Case 3: The Halted Stock

The Setup: Carol is an LP in a GameStop/USDC pool with \$20,000 of liquidity. GameStop is at \$25.

What Happens:

1. **10:30 AM:** GameStop announces a stock split. NYSE halts trading for “news pending.”
2. **10:30-10:45 AM:** No trades happen on NYSE. But Carol's AMM is technically still open (it's a blockchain—it doesn't know about halts).
3. **Problem 1:** If people can still trade on the AMM, they might:

- Trade based on incomplete information
 - Manipulate the pool price with no reference market
 - Force Carol to provide liquidity at stale prices
4. **Problem 2:** The oracle is stuck. It was getting prices from NYSE, which is now halted. The oracle shows \$25, but nobody knows what the “real” price should be.
5. **10:45 AM:** Trading resumes. GameStop opens at \$28. Everyone who bought from Carol’s pool at \$25 made instant profits.

Key Insight

The Protection: The AMM needs a “trading halt detector”:

- If the oracle hasn’t updated in 5+ minutes during market hours, assume a halt
- Automatically pause AMM trading
- Resume only when oracle shows fresh prices

4.4 Edge Case 4: The Low-Liquidity Hour Attack

The Setup: Dave is an LP in an Amazon/USDC pool. It’s 6:00 AM ET (pre-market). His pool has \$30,000 of liquidity. Amazon is at \$180.

What Happens:

1. **6:00 AM:** A whale wants to accumulate Amazon tokens cheaply. Pre-market ECN liquidity is thin (only \$500,000 on the entire pre-market book).
2. **The Strategy:** The whale:
 - (a) Places a large sell order on the ECN, pushing pre-market price to \$175
 - (b) Buys Amazon tokens from Dave’s AMM, which still shows \$178 (oracle is slow)
 - (c) Cancels their ECN order, price rebounds
 - (d) Profits from the difference
3. **Dave’s problem:** He sold Amazon tokens at \$178 when the “real” price was moving to \$175 (momentarily), then back to \$180. He lost on both legs.
4. **Why this works:** In low-liquidity hours, it’s cheap to move prices on external venues temporarily. The AMM can’t tell the difference between a real price move and manipulation.

The Protection: During pre-market/after-hours:

- Use TWAP (time-weighted average price) oracles that smooth over 5-15 minutes
- Increase spreads proportionally to oracle uncertainty
- Limit maximum trade sizes to prevent large extractions

4.5 Edge Case 5: The Flash Crash Recovery

The Setup: Emma is an LP in a Microsoft/USDC pool with \$75,000 of liquidity. Microsoft is at \$400.

What Happens:

1. **2:45 PM:** A fat-finger error causes a brief flash crash on NYSE. Microsoft drops to \$350 for about 45 seconds, then recovers to \$398.
2. **The Oracle Problem:** The oracle picks up the \$350 price and broadcasts it.
3. **Arbitrage Window:** For those 45 seconds:
 - Emma's AMM reprices to \$350
 - Arbitrageurs buy from the AMM at \$350
 - They sell on NYSE as it recovers, pocketing \$48 per share
4. **Emma's loss:** She sold at crash prices, but the crash wasn't real—it was a technical glitch that reversed in under a minute.
5. **The Double Whammy:** When the price recovers, arbitrageurs sell back into her pool, and she might lose again on the way back up.

The Numbers:

Phase	Emma's Loss
Flash crash (\$400 → \$350)	Sold MSFT cheap: -\$2,500
Recovery (\$350 → \$398)	Bought MSFT expensive: -\$2,400
Net loss from 45-second glitch	-\$4,900

Warning

Flash crashes happen more often than people think. The 2010 Flash Crash lasted 36 minutes. Mini flash crashes of 30-60 seconds happen multiple times per year on major stocks.

The Protection:

- Use median of multiple oracle sources (if one shows \$350 but two others show \$395, use \$395)
- Implement “circuit breakers” that pause trading if price moves more than 5% in under 1 minute
- Use TWAPs that smooth out momentary spikes

4.6 Edge Case 6: The Informed Trader

The Setup: Frank is an LP in a Pfizer/USDC pool with \$40,000 of liquidity. Pfizer is at \$30.

What Happens:

1. **Background:** An FDA committee is meeting to review Pfizer's new drug. The decision will be announced at 4:30 PM.
2. **4:15 PM:** Someone with inside knowledge (or very good analysis) believes the drug will be approved and Pfizer will jump 20%.

3. **The Trade:** This informed trader buys as many Pfizer tokens as possible from Frank's pool, paying \$30.50 on average (with slippage).
4. **4:30 PM:** FDA approves the drug. Pfizer jumps to \$36.
5. **Frank's loss:** He sold tokens at \$30.50 that are now worth \$36—a \$5.50 loss per token. On 500 tokens, that's \$2,750.
6. **Why Frank couldn't win:** The informed trader *knew* something Frank didn't. This is adverse selection in action. Frank was always going to be on the losing side of this trade.

Key Insight

This is the core problem described by Glosten-Milgrom: LPs systematically lose to informed traders. The only defense is to charge enough fees to offset these losses, or to detect unusual trading patterns that suggest informed trading is happening.

The Protection:

- Monitor VPIN (Volume-synchronized Probability of Informed Trading)
- If VPIN spikes before a known announcement, widen spreads dramatically
- Consider pausing trading during known high-uncertainty events (FDA decisions, earnings, etc.)

4.7 Edge Case 7: The Holiday Weekend Disaster

The Setup: Grace is an LP in a Bank of America/USDC pool with \$60,000 of liquidity. It's Thursday before a 3-day weekend (Friday is a market holiday). BofA is at \$35.

What Happens:

1. **Thursday 4:00 PM:** NYSE closes for a 3-day weekend.
2. **Friday:** A major bank in Europe fails. Markets in Europe crash 8%. Crypto markets (which are open) react negatively. Bank stocks are expected to drop significantly.
3. **Saturday-Sunday:** News continues to develop. Everyone expects BofA to open down 10-15% on Monday.
4. **The Problem:** Grace's AMM might still be accepting trades all weekend at \$35, while the "real" expected price is now \$30.
5. **Monday 9:30 AM:** BofA opens at \$29.50. But over the weekend, 200 tokens were sold into Grace's pool at an average of \$34.
6. **Grace's loss:** She bought 200 tokens at \$34 that are now worth \$29.50—a loss of \$900.

A Worse Scenario: What if the news is *positive*? Then arbitrageurs would be buying from Grace's pool all weekend at the old low price, benefiting even more.

Warning

Holiday weekends are especially dangerous because:

- 3-4 days of accumulated news creates larger gaps
- Thin trading volume makes manipulation easier
- Global events don't stop just because the NYSE is closed

The Protection: For holiday weekends:

- Increase fees to 2-5% (essentially making trading uneconomical except for urgent needs)
- Or pause trading entirely
- Publish expected reopening time so traders know when normal trading resumes

4.8 Edge Case 8: The Stale Oracle Attack

The Setup: Henry is an LP in a Nvidia/USDC pool with \$80,000 of liquidity. Nvidia is at \$800.

What Happens:

1. **2:00 PM:** The Chainlink oracle that feeds Nvidia's price has a temporary issue. It stops updating.
2. **2:00-2:15 PM:** Neural network breakthrough is announced. Nvidia jumps from \$800 to \$850 on NYSE.
3. **The AMM's Problem:** It still thinks Nvidia is \$800 because the oracle is stale.
4. **The Attack:** An attacker buys Nvidia from Henry's pool at \$800, immediately sells on NYSE at \$850. Repeats until Henry's pool is drained of Nvidia tokens.
5. **Henry's loss:** He sold \$25,000 worth of Nvidia at 6% below market price—a \$1,500 loss in 15 minutes.
6. **Why Henry didn't know:** The AMM's UI might still show “oracle updated 15 minutes ago” which seems fine normally, but during rapidly moving markets, 15 minutes is an eternity.

Key Insight

Oracle freshness requirements should be **dynamic**:

- During quiet markets: 60-second staleness is acceptable
- During volatile markets: even 10 seconds of staleness can be exploited
- The system should detect volatility and tighten staleness requirements automatically

The Protection:

- Check oracle staleness before every trade
- Block trades if oracle is more than 60 seconds old during core hours
- Use multiple oracle sources (Chainlink, Pyth, Redstone) and require 2-of-3 to be fresh
- Display prominent warnings when oracles are stale

4.9 Summary: The Common Threads

All of these edge cases share common themes:

Theme	Description
Information Asymmetry	Someone knows more than the LP (news, inside info, oracle status)
Hedging Inability	LP can't offset their risk because external markets are closed/illiquid
Oracle Lag	The AMM's view of "fair price" is stale compared to reality
Time Discontinuity	Market close → open, halt → resume, weekday → weekend
Speed Mismatch	Arbitrageurs are faster than LPs (or the protocol) can react

The protection mechanisms described in the following sections address each of these themes systematically.

5 Part 3: Risk Detection Mechanisms

To protect LPs, we first need to detect when market conditions are dangerous. This section covers the signals and metrics used to identify risky situations.

5.1 Time-Based Regime Detection

The simplest form of risk detection: know what time it is relative to traditional market hours.

5.1.1 Deterministic Market Hours Calendar

For NYSE-traded securities:

Regime	Hours (ET)	Risk Level	Characteristics
Overnight	8:00 PM – 4:00 AM	Very High	No hedge venue, thin data
Pre-Market Early	4:00 AM – 7:00 AM	High	Limited ECN liquidity
Pre-Market Late	7:00 AM – 9:30 AM	Medium-High	Building activity
Soft Open	9:30 AM – 9:35 AM	Medium	Gap resolution period
Core Session	9:35 AM – 4:00 PM	Low	Deep liquidity, tight spreads
After-Hours	4:00 PM – 8:00 PM	Medium-High	Declining liquidity
Weekend	Fri 8 PM – Mon 4 AM	Very High	No trading anywhere

5.1.2 Pseudocode for Regime Detection

Listing 1: Time-Based Regime Detection

```

1 def get_market_regime(current_time_utc):
2     # Convert to Eastern Time
3     et_hour = get_eastern_hour(current_time_utc)
4     et_minute = get_eastern_minute(current_time_utc)
5     day_of_week = get_day_of_week(current_time_utc)
6
7     # Weekend check
8     if day_of_week in [SATURDAY, SUNDAY]:

```

```

9     return REGIME_WEEKEND
10
11    # Holiday check (requires calendar oracle)
12    if is_market_holiday(current_time_utc):
13        return REGIME_HOLIDAY
14
15    # Time-based check
16    if et_hour < 4:
17        return REGIME_OVERNIGHT
18    elif et_hour < 7:
19        return REGIME_PRE_MARKET_EARLY
20    elif et_hour < 9 or (et_hour == 9 and et_minute < 30):
21        return REGIME_PRE_MARKET_LATE
22    elif et_hour == 9 and et_minute < 35:
23        return REGIME_SOFT_OPEN
24    elif et_hour < 16:
25        return REGIME_CORE_SESSION
26    elif et_hour < 20:
27        return REGIME_AFTER_HOURS
28    else:
29        return REGIME_OVERNIGHT

```

5.2 VPIN: Volume-Synchronized Probability of Informed Trading

VPIN is a real-time metric for detecting “toxic” order flow—order flow that’s more likely to be informed.

5.2.1 The Intuition

Traditional time-based metrics (like hourly volatility) don’t capture the information content of trades well. VPIN uses **volume time** instead—dividing the trading day into buckets of equal volume rather than equal time.

5.2.2 Calculation Steps

1. **Volume Bucketing:** Divide all trades into fixed-size volume buckets. For example, if you choose bucket size = 1,000 shares, then bucket 1 contains the first 1,000 shares traded, bucket 2 contains the next 1,000, and so on.
2. **Trade Classification:** For each bucket, classify how much of the volume was “buy” versus “sell”. This can be done using:
 - The **tick rule**: if the trade price is higher than the previous trade, it’s a buy
 - The **quote rule**: if the trade is closer to the ask, it’s a buy
 - A probabilistic classification based on trade size and timing
3. **Imbalance Calculation:** For each bucket i , calculate the absolute imbalance:

$$\text{Imbalance}_i = |V_i^{\text{buy}} - V_i^{\text{sell}}|$$

4. **VPIN Calculation:** Average the imbalances over the last n buckets:

$$\text{VPIN} = \frac{\sum_{i=1}^n |V_i^{\text{buy}} - V_i^{\text{sell}}|}{n \times V_{\text{bucket}}}$$

VPIN Formula:

$$\text{VPIN} = \frac{\sum_{i=1}^n |V_i^{buy} - V_i^{sell}|}{n \cdot V_{bucket}}$$

Where:

- n = number of buckets in the lookback window
- V_{bucket} = size of each volume bucket
- V_i^{buy}, V_i^{sell} = buy and sell volumes in bucket i

5.2.3 Interpreting VPIN

VPIN Level	Interpretation	Recommended Action
< 0.3	Normal conditions	Standard spreads
0.3 – 0.5	Elevated informed trading	Widen spreads 20–50%
0.5 – 0.7	High toxicity	Widen spreads 50–100%, reduce depth
> 0.7	Extreme toxicity	Consider temporary pause

Key Insight

VPIN spiked to extreme levels **hours before** the 2010 Flash Crash, providing early warning of market stress. It's one of the best real-time indicators of order flow toxicity.

5.3 Price Deviation from Oracle

A simple but effective check: how far is the AMM's price from a trusted external price?

Price Deviation:

$$\text{Deviation} = \frac{|P_{\text{AMM}} - P_{\text{oracle}}|}{P_{\text{oracle}}} \times 10000 \text{ (in basis points)}$$

Thresholds by regime:

Regime	Warning Threshold	Block Threshold
Core Session	100 bps (1%)	300 bps (3%)
Pre/Post Market	200 bps (2%)	500 bps (5%)
Overnight	300 bps (3%)	1000 bps (10%)

5.4 Oracle Staleness Check

An oracle that hasn't updated recently may be providing outdated prices.

Staleness Check:

Is Stale = (current time – oracle update time) > max staleness

Maximum staleness by regime:

Regime	Max Staleness
Core Session	60 seconds
Pre/Post Market	120 seconds
Overnight	300 seconds (or block all trades)

6 Part 4: Protection Mechanisms

Now we can design the actual protection mechanisms that respond to the detected risks.

6.1 Dynamic Fee Adjustment

Fees should increase when conditions are riskier, compensating LPs for the higher expected losses.

Dynamic Fee Formula:

$$\text{Fee}(t) = \text{BaseFee} + \alpha \cdot \sigma_t^2 + \beta \cdot \text{VPIN}_t + \gamma \cdot \text{RegimeMult}_t + \delta \cdot |\text{Inventory}_t|$$

Where:

- BaseFee = minimum fee (e.g., 0.05%)
- α = volatility sensitivity
- σ_t = current realized volatility
- β = toxicity sensitivity
- VPIN_t = current VPIN measure
- γ = regime sensitivity coefficient
- RegimeMult_t = multiplier based on market hours
- δ = inventory sensitivity
- Inventory_t = current pool imbalance

Example parameter values:

Regime	BaseFee	α	β	RegimeMult	δ
Core Hours	0.05%	0.5	0.3	1.0	0.1
Pre/Post Market	0.15%	1.0	0.5	2.0	0.2
Overnight	0.30%	1.5	1.0	4.0	0.3
Weekend	0.50%	2.0	1.5	6.0	0.5

6.2 Gap Capture Mechanism

At market open, when there's a gap between the previous close price and the new open price, a special auction can capture this value for LPs instead of letting arbitrageurs take it all.

6.2.1 The Problem

Suppose Apple closed at \$150 yesterday, but due to overnight news, it will open at \$160 today. If the AMM still prices Apple at \$150 at 9:30 AM, arbitrageurs will buy as much as possible from the AMM and sell at \$160 on the NYSE. LPs lose \$10 per share.

6.2.2 The Solution: Minimum Bid with Decay

Instead of allowing immediate trades at the stale price, require a **minimum bid** that decays over time:

Minimum Bid (Gap Capture):

$$\text{MinBid}(t) = \text{Gap} \times L \times \text{CaptureRate} \times e^{-\lambda(t-t_{\text{open}})}$$

Where:

- $\text{Gap} = |P_{\text{oracle}} - P_{\text{pool}}|$ = price difference
- L = liquidity affected (in \$)
- CaptureRate = target % for LPs (e.g., 70%)
- λ = decay constant (e.g., 0.4 per minute)
- $t - t_{\text{open}}$ = time since market opened

Example decay curve:

Time Since Open	Capture Rate
0 minutes	70%
1 minute	56%
2 minutes	45%
3 minutes	36%
4 minutes	29%
5+ minutes	0% (normal trading)

6.2.3 How It Works

1. At market open (9:30 AM), detect if there's a significant gap (e.g., $>0.5\%$).
2. If yes, enter “gap auction mode”.
3. Any trade must include a minimum bid that goes directly to LPs (on top of the normal swap).
4. The required bid decays exponentially over 5 minutes.
5. After 5 minutes, or if the gap closes due to trading, return to normal mode.

6.3 Inventory-Aware Quote Skewing

From the Avellaneda-Stoikov model, we know that quotes should be skewed based on inventory. Here's how to implement this in an AMM:

Effective Price with Inventory Skew:

$$P_{\text{effective}} = P_{\text{oracle}} \times (1 + \kappa \times \text{Imbalance})$$

Where:

- $\text{Imbalance} = \frac{\text{Value}_{\text{token0}} - \text{Value}_{\text{token1}}}{\text{Value}_{\text{token0}} + \text{Value}_{\text{token1}}}$ (using oracle prices)
- κ = inventory sensitivity (typically 0.01 to 0.05)

Effect on trading:

Imbalance	Effect (with $\kappa = 0.02$)	Interpretation
+10% (long token0)	+0.2% skew	Makes selling token0 slightly cheaper
+25% (long token0)	+0.5% skew	Strongly encourages selling token0
-10% (short token0)	-0.2% skew	Makes buying token0 slightly cheaper
-25% (short token0)	-0.5% skew	Strongly encourages buying token0

6.4 Circuit Breakers

When conditions become extreme, trading should be paused or heavily restricted.

6.4.1 Trigger Conditions

Condition	Threshold	Action
Gap at open	> 10%	Block trades until auction settles
Intraday deviation	> 5% in 5 min	Require auction certificate
Oracle staleness	> 60 seconds (core)	Block trades
Inventory imbalance	> 40%	Widen spreads 3x
VPIN	> 0.7	Emergency pause
Volume spike	> 10× average	Enter high-fee mode

6.4.2 Graduated Response

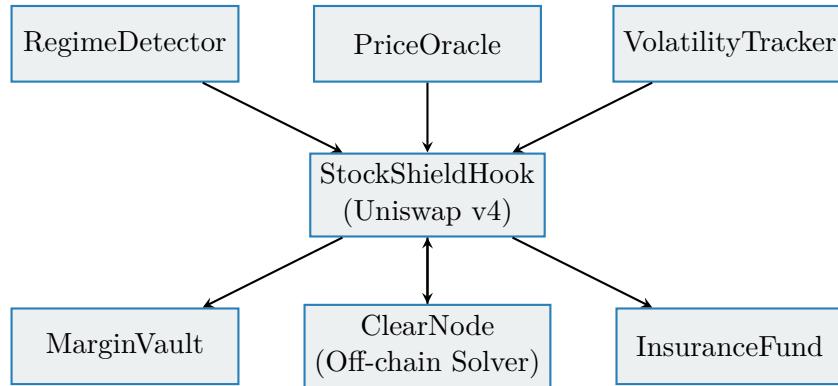
Rather than binary on/off, use graduated responses:

Level	Condition	Spread	Depth	Notes
Normal	No triggers	1x	100%	Standard operation
Level 1 (Warning)	1 minor trigger	2x	100%	Increased caution
Level 2 (Caution)	2 triggers or 1 major	5x	50%	Significant restrictions
Level 3 (Danger)	Multiple major triggers	10x	25%	Near-pause conditions
Level 4 (Pause)	Extreme conditions	∞	0%	No trading allowed

7 Part 5: Implementation Architecture

This section describes how to implement LP protection in a Uniswap v4 hook architecture.

7.1 System Components



7.2 Key Data Structures

Listing 2: Market State Structure

```

1 struct MarketState {
2     // Regime detection
3     Regime currentRegime;
4     uint256 regimeStartTime;
5
6     // Price tracking
7     uint256 lastOraclePrice;
8     uint256 lastOracleUpdate;
9     uint256 poolPrice;
10
11    // Volatility
12    uint256 realizedVolatility; // EMA of squared returns
13    uint256 lastVolatilityUpdate;
14
15    // Inventory
16    uint256 token0Value;
17    uint256 token1Value;
18    int256 imbalance; // Signed, scaled by 1e18
19
20    // VPIN (simplified)
21    uint256 recentBuyVolume;
22    uint256 recentSellVolume;
23    uint256 vpInScore; // Scaled by 1e18
24
25    // Protection state
26    uint8 circuitBreakerLevel; // 0-4
27    bool inGapAuction;
28    uint256 gapAuctionEndTime;
29 }
  
```

7.3 Hook Implementation Outline

Listing 3: Protection Hook Before Swap

```

1 function beforeSwap(
2     address sender,
3     PoolKey calldata key,
4     IPoolManager.SwapParams calldata params,
  
```

```

5     bytes calldata hookData
6 ) external override returns (bytes4, BeforeSwapDelta, uint24) {
7
8     MarketState storage state = markets[key.toId()];
9
10    // 1. Update regime
11    state.currentRegime = regimeDetector.getCurrentRegime();
12
13    // 2. Check oracle freshness
14    (uint256 oraclePrice, uint256 updateTime) = oracle.getPrice(key);
15    require(
16        block.timestamp - updateTime <= getMaxStaleness(state.
17            currentRegime),
18        "Oracle stale"
19    );
20
21    // 3. Check price deviation
22    uint256 deviation = calculateDeviation(state.poolPrice, oraclePrice
23        );
24    require(
25        deviation <= getMaxDeviation(state.currentRegime),
26        "Price deviation too high"
27    );
28
29    // 4. Check circuit breaker
30    require(
31        state.circuitBreakerLevel < LEVEL_PAUSE,
32        "Trading paused"
33    );
34
35    // 5. Calculate dynamic fee
36    uint24 dynamicFee = calculateDynamicFee(state);
37
38    // 6. If in gap auction, require minimum bid
39    if (state.inGapAuction) {
40        uint256 requiredBid = calculateGapBid(state);
41        // Decode and verify bid from hookData
42        require(getBidFromHookData(hookData) >= requiredBid, "Bid too
43        low");
44    }
45
46    // Return dynamic fee to pool
47    return (
48        BaseHook.beforeSwap.selector,
49        BeforeSwapDeltaLibrary.ZERO_DELTA,
50        dynamicFee | LPFeeLibrary.OVERRIDE_FEE_FLAG
51    );
52}

```

7.4 Volatility Tracking (On-Chain)

Listing 4: Exponential Moving Average Volatility

```

1  uint256 constant SMOOTHING_FACTOR = 20; // 20-period EMA
2
3  function updateVolatility(MarketState storage state, uint256 newPrice)
4      internal {

```

```

4     if (state.lastOraclePrice == 0) {
5         state.lastOraclePrice = newPrice;
6         return;
7     }
8
9     // Calculate log return squared (approximation)
10    // logReturn ~= (newPrice - oldPrice) / oldPrice for small changes
11    int256 priceDiff = int256(newPrice) - int256(state.lastOraclePrice)
12    ;
13    uint256 absReturn = priceDiff >= 0
14        ? uint256(priceDiff) * 1e18 / state.lastOraclePrice
15        : uint256(-priceDiff) * 1e18 / state.lastOraclePrice;
16
17    uint256 returnSquared = absReturn * absReturn / 1e18;
18
19    // EMA update: new = alpha * current + (1-alpha) * old
20    // alpha = 2 / (SMOOTHING_FACTOR + 1)
21    uint256 alpha = (2 * 1e18) / (SMOOTHING_FACTOR + 1);
22    state.realizedVolatility =
23        returnSquared * alpha +
24        state.realizedVolatility * (1e18 - alpha)
25    ) / 1e18;
26
27    state.lastOraclePrice = newPrice;
28    state.lastVolatilityUpdate = block.timestamp;
29}
30
31 function getAnnualizedVolatility(MarketState storage state)
32     internal view returns (uint256)
33 {
34     // Assume updates happen roughly every minute
35     // sqrt(variance_per_minute * 525600 minutes/year)
36     return sqrt(state.realizedVolatility * 525600);
}

```

8 Part 6: Parameter Calibration

The effectiveness of LP protection depends critically on choosing appropriate parameters.

8.1 Fee Parameters

Parameter	Suggested Value	Rationale
BaseFee (core hours)	5 bps	Competitive with centralized exchanges
BaseFee (overnight)	30 bps	Compensate for inability to hedge
α (volatility sensitivity)	0.5	LVR is $\sigma^2/8$, so fee $\propto \sigma^2$
β (VPIN sensitivity)	0.3	Informed trading costs extra
κ (inventory sensitivity)	0.02	2% skew per 100% imbalance

8.2 Circuit Breaker Thresholds

Threshold	Value	Rationale
Oracle staleness (core)	60 sec	Exchange-grade freshness
Oracle staleness (extended)	120 sec	Less frequent updates tolerable
Max deviation (core)	300 bps	Allow normal price discovery
Max deviation (extended)	500 bps	Wider bands for less liquid hours
Inventory imbalance warning	40%	Significant one-sided flow
VPIN critical threshold	0.7	Historical flash crash indicator

8.3 Gap Auction Parameters

Parameter	Value	Rationale
Capture rate	70%	LPs get majority, leave room for profit
Decay constant λ	0.4/min	Converge to normal in ~5 minutes
Minimum gap to trigger	50 bps	Don't trigger for normal noise
Maximum auction duration	5 min	Don't delay normal trading too long

9 Part 7: Hackathon Prize Track Integration

This section outlines how the StockShield LP Protection system strategically integrates technologies from three prize sponsors to create a cohesive, production-ready solution.

9.1 Technology-to-Purpose Mapping

Technology	Prize Track	Purpose in StockShield
Uniswap v4 Hooks	Privacy DeFi (\$5,000)	Core on-chain LP protection logic, dynamic fee adjustment, circuit breakers
Yellow Network	State Channels (TBD)	Off-chain order matching, instant settlement, cross-chain clearing
ENS	Decentralized Naming (\$5,000)	Human-readable LP vault identities, trader reputation, protocol registry

9.2 Uniswap v4: Privacy-Enhancing LP Protection

Uniswap v4's hook architecture is the **core on-chain layer** of StockShield. The Privacy DeFi track specifically asks for projects that:

- Reduce unnecessary information exposure during on-chain trading
- Improve execution quality
- Design market behavior resilient to adverse selection and extractive dynamics

9.2.1 How StockShield Addresses Privacy DeFi Requirements

Requirement	StockShield Implementation
Reduce information exposure	Gap auction bids are sealed, preventing front-running. Intent-based trading hides order direction until execution.
Improve execution quality	Dynamic fees ensure LPs are fairly compensated, reducing adverse selection. Price skewing directs flow to balanced pools.
Resilience to extraction	VPIN detection identifies toxic flow. Circuit breakers halt trading during manipulation attempts.

9.2.2 Uniswap v4 Hook Implementation Points

Listing 5: StockShield Hook Entry Points

```

1 // Hook permissions for LP Protection
2 function getHookPermissions() public pure returns (Hooks.Permissions
3     memory) {
4     return Hooks.Permissions({
5         beforeSwap: true,           // Dynamic fees, circuit breakers
6         afterSwap: true,           // Update VPIN, inventory tracking
7         beforeAddLiquidity: true,   // Gap auction participation
8         afterAddLiquidity: true,    // LP rewards distribution
9         beforeRemoveLiquidity: true, // Time-lock during high-risk
10        periods
11        afterRemoveLiquidity: true,
12        beforeDonate: false,
13        afterDonate: false
14    });
}

```

Key Hook Functions:

1. **beforeSwap**: Implements dynamic fee calculation based on regime, VPIN, and inventory. Returns the adjusted fee to the pool manager.
2. **afterSwap**: Updates on-chain state including buy/sell volume buckets for VPIN, inventory imbalance, and realized volatility EMA.
3. **beforeAddLiquidity**: During gap auctions, verifies that LPs are participating in the auction mechanism to capture gap value.
4. **beforeRemoveLiquidity**: Implements time-locks during high-risk periods (overnight, weekends) to prevent LP flight before expected gaps.

Key Insight

Privacy Enhancement: The hook's gap auction uses a commit-reveal scheme. Traders commit to a hash of their bid in Phase 1, then reveal in Phase 2. This prevents MEV bots from front-running the auction by ensuring bid amounts are only known after the commitment deadline.

9.3 Yellow Network: Off-Chain Scaling via State Channels

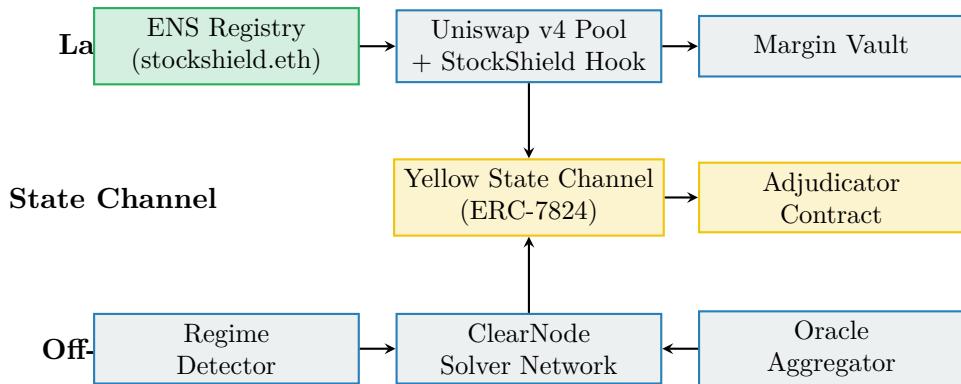
Yellow Network's state channel infrastructure (ERC-7824) enables **instant, gasless, cross-chain** order matching while maintaining full non-custody.

9.3.1 Why State Channels for LP Protection?

State Channel Benefit	LP Protection Application
Instant Finality	LPs receive protection decisions in milliseconds, not 12+ seconds. Critical for fast-moving markets.
Zero Gas Fees	High-frequency VPIN updates and inventory rebalancing become economically viable.
Cross-Chain	LP positions can be protected across multiple chains simultaneously.
Non-Custodial	LPs retain full control via on-chain margin accounts with off-chain matching.

9.3.2 Integration Architecture

The system uses a hybrid on-chain/off-chain architecture:



9.3.3 ERC-7824 State Channel Flow

- Channel Setup:** LPs open a state channel by locking collateral in the on-chain Margin Vault. The channel is identified by a unique channel ID linked to their ENS name.
- Off-Chain Updates:** During normal operation, VPIN calculations, inventory updates, and fee adjustments happen entirely off-chain in the ClearNode network. Each update is cryptographically signed by both parties.
- Periodic Settlement:** Every N blocks (configurable), the latest signed state is submitted on-chain to the Uniswap v4 hook, updating the pool's fee tier and protection parameters.
- Dispute Resolution:** If parties disagree, the Adjudicator contract accepts the latest mutually-signed state. The on-chain fallback ensures safety even if the state channel network is unavailable.

Listing 6: Yellow State Channel State Structure

```

1 struct ChannelState {
2     // Channel identification
3     bytes32 channelId;
4     address lpAddress;
5     bytes32 ensNode;           // ENS identifier for human-readable naming
6

```

```

7   // Protection parameters (updated off-chain)
8   uint256 currentVPIN;           // Latest VPIN score (1e18 scaled)
9   Regime currentRegime;         // Market regime classification
10  uint24 recommendedFee;        // Fee based on off-chain calculations
11  int256 inventory;            // Current inventory imbalance
12
13  // State channel mechanics
14  uint256 turnNum;             // Monotonically increasing state number
15  uint256 timestamp;
16  bytes32 stateHash;          // Hash of full state for verification
17
18  // Signatures
19  bytes lpSignature;
20  bytes clearNodeSignature;
21 }
```

Key Insight

Yellow Integration Value: By moving VPIN calculation and regime detection off-chain, StockShield can update protection parameters 100x+ faster than on-chain alternatives. The state channel ensures these updates remain non-custodial and verifiable.

9.4 ENS: Decentralized Identity Layer

ENS provides the **human-readable identity layer** for StockShield, enabling:

- Human-readable LP vault names (e.g., `alice-lp.stockshield.eth`)
- Trader reputation linking across protocols
- Protocol parameter registry
- Cross-chain identity resolution

9.4.1 ENS Use Cases in StockShield

Use Case	Implementation
LP Vault Identity	Each LP vault gets a subname (e.g., <code>vault-42.stockshield.eth</code>). Users can deposit to vaults by name instead of address.
Trader Reputation	Traders' historical VPIN contribution (toxic vs. benign flow) is stored against their ENS name, enabling reputation-based fee tiers.
Protocol Registry	<code>registry.stockshield.eth</code> stores parameter configurations, oracle addresses, and ClearNode endpoints.
Multi-Chain Resolution	ENS CCIP-Read enables the same vault name to resolve to different addresses on different chains.

9.4.2 ENS Implementation

Listing 7: StockShield ENS Resolver

```
1 // SPDX-License-Identifier: MIT
```

```

2 pragma solidity ^0.8.24;
3
4 import "@ensdomains/ens-contracts/contracts/resolvers/Resolver.sol";
5
6 contract StockShieldResolver is Resolver {
7     // LP vault address resolution
8     mapping(bytes32 => address) public vaultAddresses;
9
10    // Trader reputation scores (scaled 1e18)
11    mapping(bytes32 => uint256) public reputationScores;
12
13    // Protocol parameter records
14    mapping(bytes32 => bytes) public protocolParams;
15
16    // Register a new LP vault with ENS subname
17    function registerVault(
18        string calldata subname,
19        address vaultAddress
20    ) external {
21        bytes32 node = keccak256(abi.encodePacked(
22            StockShield_NODE,
23            keccak256(bytes(subname))
24        ));
25
26        // Only vault owner can register
27        require(IVault(vaultAddress).owner() == msg.sender, "Not owner");
28
29        vaultAddresses[node] = vaultAddress;
30        emit VaultRegistered(subname, vaultAddress);
31    }
32
33    // Update trader reputation based on trading history
34    function updateReputation(
35        bytes32 ensNode,
36        uint256 vpinContribution,
37        bool wasToxic
38    ) external onlyClearNode {
39        uint256 current = reputationScores[ensNode];
40
41        if (wasToxic) {
42            // Toxic flow decreases reputation
43            reputationScores[ensNode] = (current * 95) / 100;
44        } else {
45            // Benign flow increases reputation
46            reputationScores[ensNode] = min(
47                (current * 101) / 100,
48                1e18 // Max reputation = 1.0
49            );
50        }
51    }
52}

```

9.4.3 Reputation-Based Fee Tiers

ENS enables a novel fee structure based on trader reputation:

Reputation Score	Tier	Fee Multiplier	Description
≥ 0.8	Platinum	0.8x	Consistently benign flow, minimal LP risk
0.6 – 0.8	Gold	0.9x	Mostly positive trading history
0.4 – 0.6	Silver	1.0x	Standard tier, neutral history
0.2 – 0.4	Bronze	1.2x	Some toxic flow detected
< 0.2	Restricted	2.0x	High-frequency toxic flow

Warning

Reputation must be non-gameable. The system uses a slow-moving EMA with asymmetric updates: reputation decreases 5% per toxic trade but only increases 1% per benign trade. This prevents wash trading to build reputation before a large extraction.

9.5 Integrated System: Full Stack Architecture

Layer	Technology	Components
Identity	ENS	Vault names, trader reputation, protocol registry
On-Chain	Uniswap v4	Pool contracts, AURA hook, margin vault
State Channel	Yellow (ERC-7824)	Fast settlement, cross-chain clearing
Off-Chain	ClearNode Network	VPIN calculation, regime detection, oracle aggregation

9.5.1 Complete Transaction Flow

1. Trader Initiation:

- Trader resolves `eth-usdc.stockshield.eth` via ENS to find the pool
- ENS lookup also returns trader's reputation score

2. Off-Chain Pre-Processing (Yellow):

- Trade intent sent to ClearNode via state channel
- ClearNode calculates current VPIN, regime, and recommended fee
- If multiple chains have liquidity, ClearNode finds optimal routing

3. On-Chain Execution (Uniswap v4):

- Trade submitted to Uniswap v4 with hook data containing:
 - Signed state from Yellow channel (proving current protection params)
 - Trader's ENS-linked reputation score
 - Gap auction bid (if applicable)
- Hook validates state, applies dynamic fee, executes swap

4. Post-Trade Updates:

- Hook emits events for VPIN contribution
- ClearNode updates trader reputation in ENS resolver
- State channel state is updated with new inventory

9.6 Prize Track Qualification Matrix

Track	Requirements	StockShield Deliverables	Status
Uniswap v4 Privacy DeFi	Functional code, TxID, GitHub, README, demo video	<ul style="list-style-type: none"> • Hook contract (Solidity) • Testnet deployments • 3-min demo video 	Required
Yellow Network	State channel integration, cross-chain settlement	<ul style="list-style-type: none"> • ERC-7824 channel impl • ClearNode integration • Cross-chain demo 	TBD
ENS	Decentralized naming integration	<ul style="list-style-type: none"> • StockShield Resolver contract • Vault naming system • Reputation registry 	Required

Key Insight

The three technologies are synergistic: ENS provides identity, Yellow provides speed, and Uniswap v4 provides the core financial logic. Together they create a system that is **human-readable** (ENS), **fast** (Yellow), and **secure** (Uniswap v4).

10 Conclusion

Protecting liquidity providers in tokenized securities AMMs requires a multi-layered approach:

1. **Understand the fundamentals:** LPs face both inventory risk (Avellaneda-Stoikov) and adverse selection risk (Kyle, Glosten-Milgrom). LVR ($\sigma^2/8$) quantifies the unavoidable cost of passive market making.
2. **Detect risk in real-time:** Time-based regime detection handles the deterministic aspects (market hours). VPIN and price deviation detect real-time toxicity.
3. **Respond dynamically:** Dynamic fees adjust to conditions. Gap auctions capture overnight value. Inventory skewing manages one-sided flow. Circuit breakers halt trading in extreme conditions.
4. **Implement robustly:** Uniswap v4 hooks provide the flexibility needed for custom protection logic. On-chain state tracking enables autonomous operation.

Key Insight

The three risk surfaces—open/close gaps, session transitions, and low-liquidity hours—are all manifestations of one core problem: **discontinuous hedging ability**. A unified protection framework that detects regimes and adjusts pricing accordingly is the solution.

References

1. Avellaneda, M., & Stoikov, S. (2008). High-frequency trading in a limit order book. *Quantitative Finance*, 8(3), 217-224.

2. Guéant, O., Lehalle, C. A., & Fernandez-Tapia, J. (2013). Dealing with the inventory risk: A solution to the market making problem. *Mathematics and Financial Economics*, 7(4), 477-507.
3. Kyle, A. S. (1985). Continuous auctions and insider trading. *Econometrica*, 53(6), 1315-1335.
4. Glosten, L. R., & Milgrom, P. R. (1985). Bid, ask and transaction prices in a specialist market with heterogeneously informed traders. *Journal of Financial Economics*, 14(1), 71-100.
5. Easley, D., López de Prado, M. M., & O'Hara, M. (2012). Flow toxicity and liquidity in a high-frequency world. *The Review of Financial Studies*, 25(5), 1457-1493.
6. Milionis, J., Moallemi, C. C., Roughgarden, T., & Zhang, A. L. (2023). Automated market making and loss-versus-rebalancing. *arXiv preprint arXiv:2208.06046*.
7. Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2), 357-384.
8. Hendershott, T., & Seasholes, M. S. (2007). Market maker inventories and stock prices. *American Economic Review*, 97(2), 210-214.
9. Stoll, H. R. (1989). Inferring the components of the bid-ask spread: Theory and empirical tests. *The Journal of Finance*, 44(1), 115-134.
10. Paradigm Research. (2021). TWAMM: Time-Weighted Average Market Maker.